

(3) [25pts] This task is inferential. You are going to try to explain why a constructor wins a season between 1950 and 2010. Fit a model using features that make theoretical sense to describe F1 racing between 1950 and 2010. Clean the data, and transform it as necessary, including dealing with missing data. [Remember, this will almost necessarily be an overfit model where variables are selected because they make sense to explain F1 races between 1950 and 2010, and not based on algorithmic feature selection]

From your fitted model:

- describe your model, and explain why each feature was selected
- provide statistics that show how well the model fits the data
- what is the most important variable in your model? How did you determine that?
- provide some marginal effects for the variable that you identified as the most important in the model, and interpret it in the context of F1 races: in other words, give us the story that the data is providing you about constructors that win seasons
- does it make sense to think of it as an "explanation" for why a constructor wins a season? or is it simply an association we observe in the data?

## Data Preprocessing:

1. Merge preprocessed data that describe features of a race with the status of each race. The status is a binary variable, where if wins equals to 1, the driver wins the race.
2. Calculated the average age of the driver for each constructor in each season.
3. Created a season variable to account for the role of seasonality in F1 race by extracting and grouping date of the race. One-hot-encode the four seasons to fit the model.
4. Created a continent variable to account for constructor's nationality. This is a continuous variable where 1,2,3, and 4 each means if the constructor is from Europe, North and South America, Asia, and Africa.
5. Grouped by year and constructor ID and Calculated the average of each explanatory variable for every constructor in every year between 1950 and 2010. Calculated the average number of wins for each constructor in each season

## Model Building:

1. Run a stats model Linear Regression model to explain why a constructor wins a season by looking at the variable's coefficients and statistical significance.
2. Build a machine learning linear regression model to look at model evaluation metrics and determine if our model is overfitting
3. Feature selection: race\_year, points, constructorId, grid, age\_as\_of\_race, positionOrder, points\_x, s\_autumn, s\_spring, s\_summer, s\_winter.
4. Split dataset into training and test set
5. Fit linear regression models
6. Find model stats (**rmse**), select the best model, and most important feature ("**points**")

## Model Result & Interpretation:

-- Stats model Result Table --

| OLS Regression Results |                  |                     |           |       |        |        |
|------------------------|------------------|---------------------|-----------|-------|--------|--------|
| Dep. Variable:         | wins             | R-squared:          | 0.609     |       |        |        |
| Model:                 | OLS              | Adj. R-squared:     | 0.605     |       |        |        |
| Method:                | Least Squares    | F-statistic:        | 156.3     |       |        |        |
| Date:                  | Sat, 10 Apr 2021 | Prob (F-statistic): | 2.60e-177 |       |        |        |
| Time:                  | 00:12:59         | Log-Likelihood:     | -355.12   |       |        |        |
| No. Observations:      | 912              | AIC:                | 730.2     |       |        |        |
| Df Residuals:          | 902              | BIC:                | 778.4     |       |        |        |
| Df Model:              | 9                |                     |           |       |        |        |
| Covariance Type:       | nonrobust        |                     |           |       |        |        |
|                        | coef             | std err             | t         | P> t  | [0.025 | 0.975] |
| Intercept              | -0.0357          | 0.085               | -0.419    | 0.676 | -0.203 | 0.132  |
| grid                   | -0.0063          | 0.003               | -2.358    | 0.019 | -0.011 | -0.001 |
| positionOrder          | 0.0084           | 0.004               | 2.049     | 0.041 | 0.000  | 0.016  |
| points_x               | 0.3000           | 0.012               | 25.998    | 0.000 | 0.277  | 0.323  |
| laps                   | -0.0006          | 0.001               | -1.158    | 0.247 | -0.002 | 0.000  |
| age_as_of_race         | 0.0025           | 0.003               | 0.867     | 0.386 | -0.003 | 0.008  |
| s_autumn               | 0.0517           | 0.066               | 0.787     | 0.431 | -0.077 | 0.181  |
| s_spring               | -0.0363          | 0.060               | -0.606    | 0.545 | -0.154 | 0.081  |
| s_summer               | -0.0092          | 0.048               | -0.193    | 0.847 | -0.103 | 0.084  |
| s_winter               | -0.0419          | 0.064               | -0.653    | 0.514 | -0.168 | 0.084  |
| continent              | -0.0577          | 0.034               | -1.722    | 0.085 | -0.124 | 0.008  |
| Omnibus:               | 269.754          | Durbin-Watson:      | 1.976     |       |        |        |
| Prob(Omnibus):         | 0.000            | Jarque-Bera (JB):   | 4366.285  |       |        |        |
| Skew:                  | 0.897            | Prob(JB):           | 0.00      |       |        |        |
| Kurtosis:              | 13.568           | Cond. No.           | 1.86e+17  |       |        |        |

### Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 1.38e-28. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

- **R-squared** is 0.609, indicating the linear regression accounts for **61%** of the variance. The model fits the data somewhat well.

- **Grid, PositionOrder**, and **points** are all **statistically significant** features. Points have the highest coefficient, which makes sense given it is used to determine the outcome of both the Drivers' and Constructor's World Championships. Each season, the Championship is awarded to the driver and the team that has scored the highest number of points.
- **Interpretation:** Holding all other variables at fixed values, we will see a **30%** increase in the **odds** of the constructor winning the season for every extra **point** the constructor earns. Points are awarded based on the position the driver finishes at.

-- Machine Learning model Result Table --

```
rmse: 0.545154665782234
rmse_train: 0.5716539710859395
mae: 0.338627292372811
R2: -4.172972188154489e-05
```

- Comparing the RMSE score on the training set and the test set, we can see that the RMSE score fitted on the training set is higher than that fitted on the test set. Therefore, our model is overfitting as expected.

Date: 2021-04-17 22:48:31

Source:  Q3\_Inference\_Model

Duration: 4.3s

Status: FINISHED

▼ Notes 

None

▼ Parameters

| Name     | Value   |
|----------|---------|
| alpha    | 43771.0 |
| l1_ratio | 0.0     |

▼ Metrics

| Name   | Value     |
|--|-----------|
| <a href="#">mae</a>   | 0.339     |
| <a href="#">r2</a>    | -4.173e-5 |
| <a href="#">rmse</a>  | 0.545     |

Showing 46 matching runs

|                          |  |        |          |              |             |                                   |         |                                |            |  |       |           |           |                    |  |       |  |
|--------------------------|--|--------|----------|--------------|-------------|-----------------------------------|---------|--------------------------------|------------|--|-------|-----------|-----------|--------------------|--|-------|--|
| Compare                  |  | Delete |          | Download CSV |             | <div></div> <div></div>           |         | Columns                        |            | <div></div> metrics.rmse < 1 and params.model = "tree" |       | Search    |           | <div></div> Filter |  | Clear |  |
|                          |  |        |          |              |             |                                   |         |                                | Parameters |  |       | Metrics > |           |                    |  |       |  |
| <input type="checkbox"/> | Start Time                                 |        | Run Name |              | User        | Source                            | Version | Models                         | alpha      | l1_ratio   | mae   | mse       | r2        |                    |  |       |  |
| <input type="checkbox"/> | <div><div></div> 2021-04-17 22:48:31</div> |        | -        |              | xl2673@c... | <div><div></div> Q3_inferen</div> | -       | <div><div></div> sklearn</div> | 43771.0    | 0.0  | 0.339 | -         | -4.173e-5 |                    |  |       |  |
| <input type="checkbox"/> | <div><div></div> 2021-04-17 22:48:08</div> |        | -        |              | xl2673@c... | <div><div></div> Q3_inferen</div> | -       | -                              | 43771.0    | 0.0  | -     | -         | -         |                    |  |       |  |
| <input type="checkbox"/> | <div><div></div> 2021-04-17 22:48:04</div> |        | -        |              | xl2673@c... | <div><div></div> Q3_inferen</div> | -       | <div><div></div> sklearn</div> | 43771.0    | 0.0  | 0.338 | -         | 6.940e-4  |                    |  |       |  |

Overall, by looking at the inference model, it's clear that only points, grid, and PositionOrder are statistically significant variables at explaining why a constructor might win a season. And by looking at the coefficient, we can find the total points earned by the constructor in each season has the highest explanatory power on the race result.

Furthermore, by looking at the linear regression machine learning model, we can clearly detect the overfitting issue in our model, because the RMSE score fitted on our training set is slightly higher than that fitted on the test set. This is expected because we want to include all variables that help to explain why a constructor wins a season, which means the model tries to explain more idiosyncrasies in the F1 dataset.